

```
In [1]: #import warnings
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: #import the libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [3]: # Import required libraries for Ridge, Lasso and GridSearchCV
from sklearn.linear_model import Ridge, Lasso
from sklearn.model_selection import GridSearchCV
```

```
In [4]: # Setting option to display all the columns and rows in dataset
pd.set_option('display.max_columns', None)
pd.set_option('display.max_rows', None)
```

```
In [5]: #Loading the dataset
df=pd.read_csv('train.csv')
```

```
In [6]: df.head()
```

Out[6]:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandCon
--	----	------------	----------	-------------	---------	--------	-------	----------	---------


0	1	60	RL	65.0	8450	Pave	NaN	Reg	
---	---	----	----	------	------	------	-----	-----	--

1	2	20	RL	80.0	9600	Pave	NaN	Reg	
---	---	----	----	------	------	------	-----	-----	--

2	3	60	RL	68.0	11250	Pave	NaN	IR1	
---	---	----	----	------	-------	------	-----	-----	--

3	4	70	RL	60.0	9550	Pave	NaN	IR1	
---	---	----	----	------	------	------	-----	-----	--

4	5	60	RL	84.0	14260	Pave	NaN	IR1	
---	---	----	----	------	-------	------	-----	-----	--



```
In [7]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
Data columns (total 81 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Id                    1460 non-null   int64
1   MSSubClass            1460 non-null   int64
2   MSZoning              1460 non-null   object
3   LotFrontage          1201 non-null   float64
4   LotArea              1460 non-null   int64
5   Street               1460 non-null   object
6   Alley               91 non-null     object
7   LotShape             1460 non-null   object
8   LandContour          1460 non-null   object
9   Utilities            1460 non-null   object
10  LotConfig            1460 non-null   object
11  LandSlope            1460 non-null   object
12  Neighborhood         1460 non-null   object
13  Condition1           1460 non-null   object
14  Condition2           1460 non-null   object
15  BldgType             1460 non-null   object
16  HouseStyle           1460 non-null   object
17  OverallQual          1460 non-null   int64
18  OverallCond          1460 non-null   int64
19  YearBuilt            1460 non-null   int64
20  YearRemodAdd         1460 non-null   int64
21  RoofStyle            1460 non-null   object
22  RoofMatl            1460 non-null   object
23  Exterior1st          1460 non-null   object
24  Exterior2nd          1460 non-null   object
25  MasVnrType           588 non-null    object
26  MasVnrArea           1452 non-null   float64
27  ExterQual            1460 non-null   object
28  ExterCond            1460 non-null   object
29  Foundation           1460 non-null   object
30  BsmtQual             1423 non-null   object
31  BsmtCond             1423 non-null   object
32  BsmtExposure         1422 non-null   object
33  BsmtFinType1         1423 non-null   object
34  BsmtFinSF1           1460 non-null   int64
35  BsmtFinType2         1422 non-null   object
36  BsmtFinSF2           1460 non-null   int64
37  BsmtUnfSF            1460 non-null   int64
38  TotalBsmtSF          1460 non-null   int64
39  Heating              1460 non-null   object
40  HeatingQC            1460 non-null   object
41  CentralAir           1460 non-null   object
42  Electrical            1459 non-null   object
43  1stFlrSF             1460 non-null   int64
44  2ndFlrSF             1460 non-null   int64
45  LowQualFinSF         1460 non-null   int64
46  GrLivArea            1460 non-null   int64
47  BsmtFullBath         1460 non-null   int64
48  BsmtHalfBath         1460 non-null   int64
49  FullBath             1460 non-null   int64
50  HalfBath             1460 non-null   int64
51  BedroomAbvGr        1460 non-null   int64
52  KitchenAbvGr         1460 non-null   int64
53  KitchenQual          1460 non-null   object
54  TotRmsAbvGrd        1460 non-null   int64

```

```
55 Functional      1460 non-null object
56 Fireplaces      1460 non-null int64
57 FireplaceQu     770 non-null object
58 GarageType      1379 non-null object
59 GarageYrBlt     1379 non-null float64
60 GarageFinish    1379 non-null object
61 GarageCars      1460 non-null int64
62 GarageArea      1460 non-null int64
63 GarageQual      1379 non-null object
64 GarageCond      1379 non-null object
65 PavedDrive      1460 non-null object
66 WoodDeckSF      1460 non-null int64
67 OpenPorchSF     1460 non-null int64
68 EnclosedPorch   1460 non-null int64
69 3SsnPorch       1460 non-null int64
70 ScreenPorch     1460 non-null int64
71 PoolArea        1460 non-null int64
72 PoolQC          7 non-null object
73 Fence           281 non-null object
74 MiscFeature     54 non-null object
75 MiscVal         1460 non-null int64
76 MoSold          1460 non-null int64
77 YrSold          1460 non-null int64
78 SaleType        1460 non-null object
79 SaleCondition   1460 non-null object
80 SalePrice       1460 non-null int64
dtypes: float64(3), int64(35), object(43)
memory usage: 924.0+ KB
```

```
In [8]: df.isnull().sum()
```

```

Out[8]: Id                0
        MSSubClass        0
        MSZoning          0
        LotFrontage      259
        LotArea          0
        Street           0
        Alley            1369
        LotShape         0
        LandContour      0
        Utilities        0
        LotConfig        0
        LandSlope        0
        Neighborhood     0
        Condition1       0
        Condition2       0
        BldgType         0
        HouseStyle       0
        OverallQual      0
        OverallCond      0
        YearBuilt        0
        YearRemodAdd     0
        RoofStyle        0
        RoofMatl         0
        Exterior1st      0
        Exterior2nd      0
        MasVnrType       872
        MasVnrArea       8
        ExterQual        0
        ExterCond        0
        Foundation       0
        BsmtQual         37
        BsmtCond         37
        BsmtExposure     38
        BsmtFinType1     37
        BsmtFinSF1       0
        BsmtFinType2     38
        BsmtFinSF2       0
        BsmtUnfSF        0
        TotalBsmtSF      0
        Heating          0
        HeatingQC        0
        CentralAir       0
        Electrical       1
        1stFlrSF         0
        2ndFlrSF         0
        LowQualFinSF     0
        GrLivArea        0
        BsmtFullBath     0
        BsmtHalfBath     0
        FullBath         0
        HalfBath         0
        BedroomAbvGr     0
        KitchenAbvGr     0
        KitchenQual      0
        TotRmsAbvGrd     0
        Functional       0
        Fireplaces       0
        FireplaceQu      690
        GarageType       81
        GarageYrBlt      81

```

```

GarageFinish      81
GarageCars        0
GarageArea        0
GarageQual        81
GarageCond        81
PavedDrive        0
WoodDeckSF        0
OpenPorchSF       0
EnclosedPorch     0
3SsnPorch         0
ScreenPorch       0
PoolArea          0
PoolQC            1453
Fence             1179
MiscFeature       1406
MiscVal           0
MoSold            0
YrSold            0
SaleType          0
SaleCondition     0
SalePrice         0
dtype: int64

```

```
In [9]: df.isnull().any().sum()
```

```
Out[9]: 19
```

```
In [10]: null_columns=df.columns[df.isnull().any()]
null_columns
```

```
Out[10]: Index(['LotFrontage', 'Alley', 'MasVnrType', 'MasVnrArea', 'BsmtQual',
               'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinType2',
               'Electrical', 'FireplaceQu', 'GarageType', 'GarageYrBlt',
               'GarageFinish', 'GarageQual', 'GarageCond', 'PoolQC', 'Fence',
               'MiscFeature'],
              dtype='object')
```

```
In [11]: null_columns.shape
```

```
Out[11]: (19,)
```

```
In [12]: null_count=df[null_columns].isnull().sum().sort_values(ascending=False)
null_per=(df[null_columns].isnull().sum()*100/df.shape[0]).sort_values(ascending
```

```
In [13]: null_data=pd.concat([null_count,null_per],axis=1,keys=['Count', 'Percentage'])
```

```
In [14]: null_data
```

Out[14]:

	Count	Percentage
PoolQC	1453	99.520548
MiscFeature	1406	96.301370
Alley	1369	93.767123
Fence	1179	80.753425
MasVnrType	872	59.726027
FireplaceQu	690	47.260274
LotFrontage	259	17.739726
GarageType	81	5.547945
GarageYrBlt	81	5.547945
GarageFinish	81	5.547945
GarageQual	81	5.547945
GarageCond	81	5.547945
BsmtFinType2	38	2.602740
BsmtExposure	38	2.602740
BsmtFinType1	37	2.534247
BsmtCond	37	2.534247
BsmtQual	37	2.534247
MasVnrArea	8	0.547945
Electrical	1	0.068493

```
In [15]: df.drop(columns=null_data[ null_data['Percentage'] > 15].index, inplace=True)
```

```
In [16]: df.shape
```

```
Out[16]: (1460, 74)
```

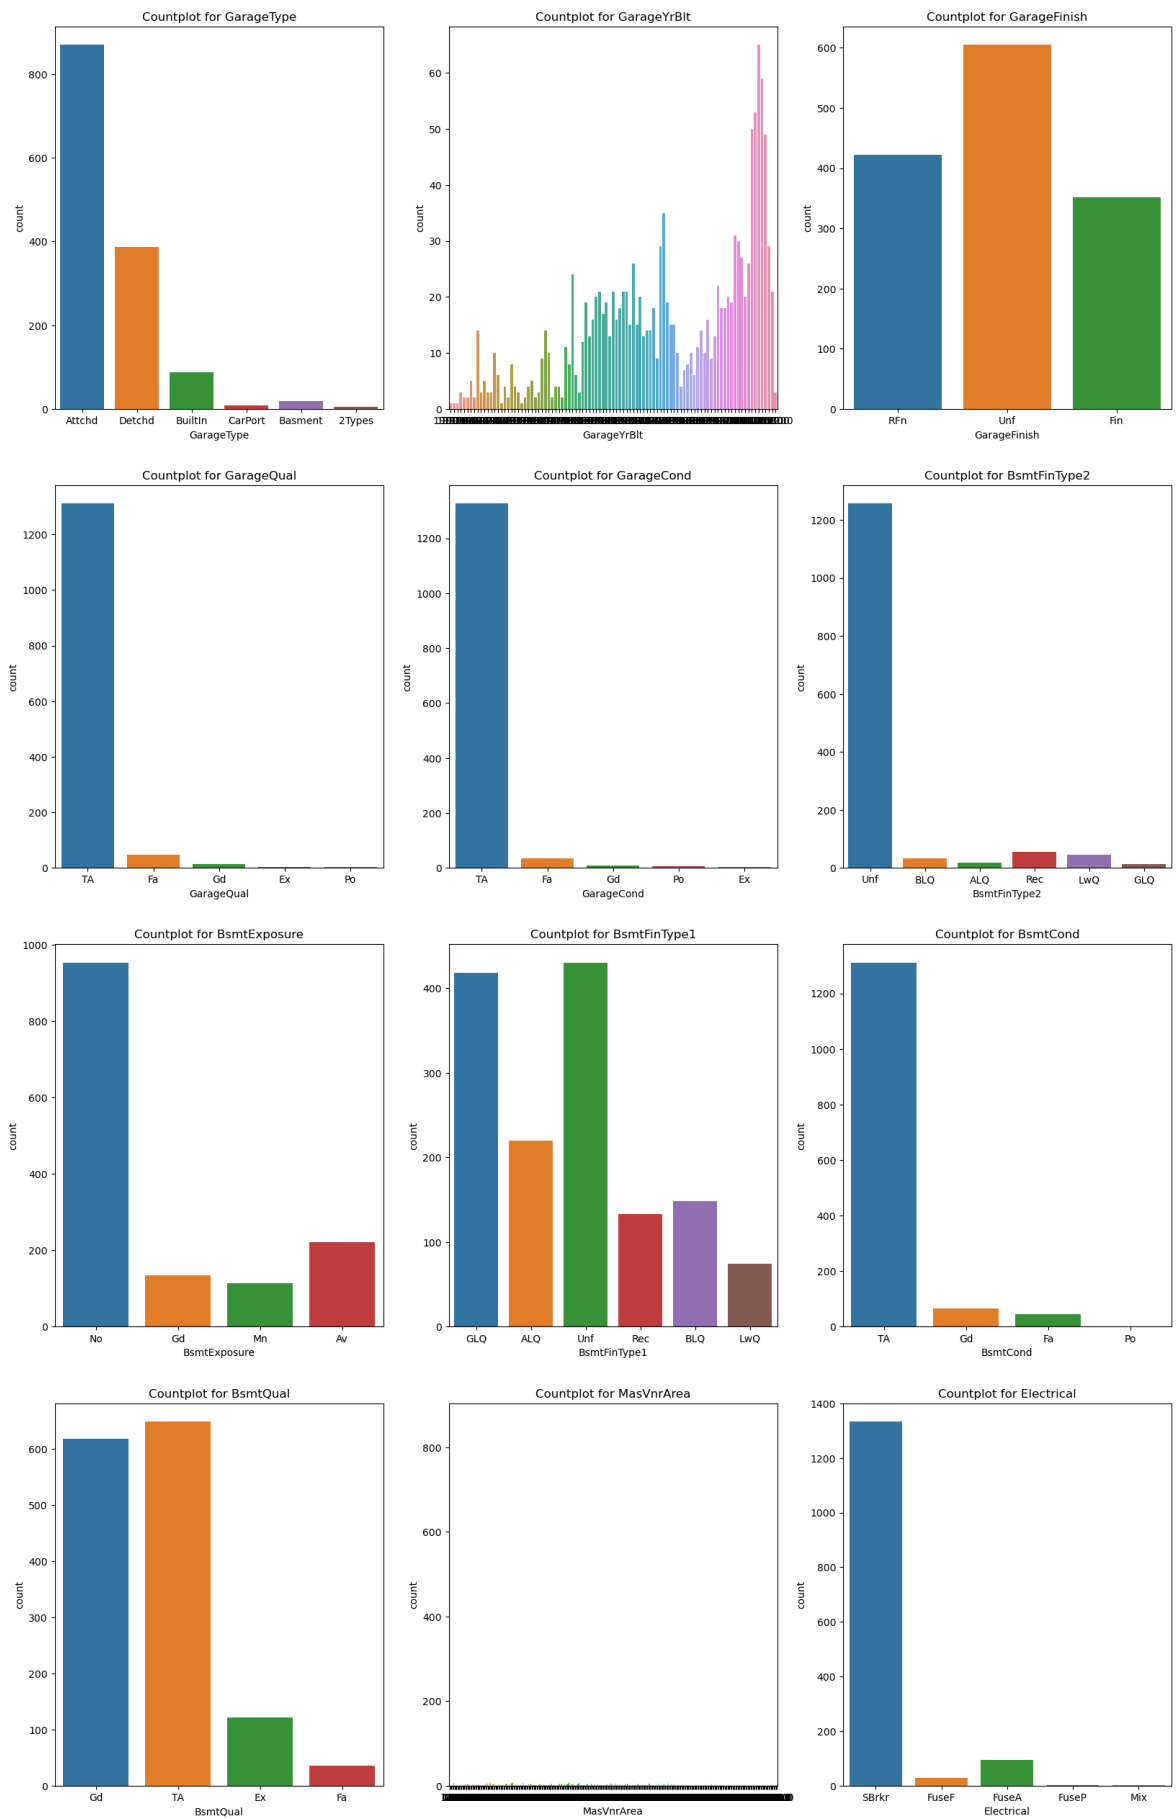
-->Use a countplot to determine which values in the columns are the most frequent.

```
In [17]: null_data = null_data[null_data['Percentage'] < 15]
null_data
```

Out[17]:

	Count	Percentage
GarageType	81	5.547945
GarageYrBlt	81	5.547945
GarageFinish	81	5.547945
GarageQual	81	5.547945
GarageCond	81	5.547945
BsmtFinType2	38	2.602740
BsmtExposure	38	2.602740
BsmtFinType1	37	2.534247
BsmtCond	37	2.534247
BsmtQual	37	2.534247
MasVnrArea	8	0.547945
Electrical	1	0.068493

```
In [18]: %matplotlib inline
plt.rcParams['figure.figsize']=20,40
for i,var in enumerate(null_data.index,start=1):
    plt.subplot(5,3,i)
    sns.countplot(x=var,data=df)
    plt.title(f"Countplot for {var}")
```



We will use the following method to conduct imputation for these columns:

- If the column is `categorical`, the missing values will be replaced using `mode()`.
- If the column contains `numerical values`, the missing values will be replaced using `median()`.

- If the value `NA` in the column has a meaningful value (GarageType = NA, for example, denotes "No Garage"). These settings will be changed to `None`.

```
In [19]: #meaningfull categorical columns
meaning_col=['GarageType','GarageFinish','GarageQual','GarageCond','BsmtExposure']
for feature in meaning_col:
    df[feature].fillna('None',inplace=True)
```

```
In [20]: df['GarageYrBlt'].fillna(df['GarageYrBlt'].median(), inplace=True)
```

```
In [21]: df['MasVnrArea'].fillna(df['MasVnrArea'].median(), inplace=True)
```

```
In [22]: df['Electrical'].fillna(df['Electrical'].mode()[0], inplace=True)
```

```
In [23]: df.isnull().any().sum()
```

```
Out[23]: 0
```

Dropping unimportant columns

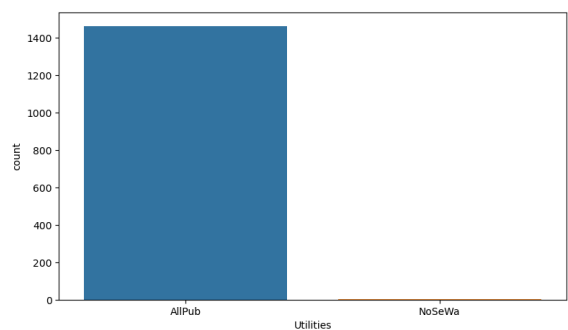






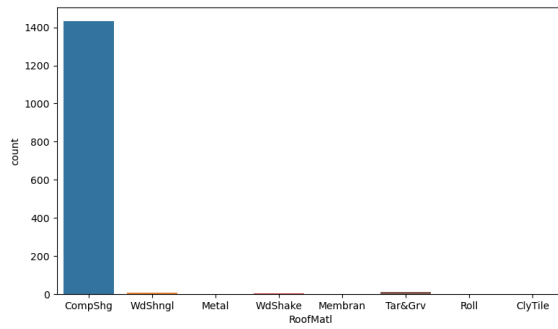


The columns will be removed in accordance with:

1. The model has less room to learn if the variance in the column is really low. These columns are going away.
2. A few columns, such as `Id`, are meaningless because they don't offer any insightful information. We'll get rid of these too

```
In [24]: df.shape
```

```
Out[24]: (1460, 74)
```

```
In [25]: %matplotlib inline
plt.rcParams['figure.figsize']=20,30
#unimportant columns
unin_col=['Street','Condition2','RoofMat1','Heating','LowQualFinSF','3SsnPorch',]
for i,var in enumerate(unin_col,start=1):
    plt.subplot(5,2,i)
    sns.countplot(x=var,data=df)
    df.drop(columns=var,inplace=True)
```



```
In [26]: df.drop(columns='Id',inplace=True)
```

```
In [27]: df.shape
```

```
Out[27]: (1460, 63)
```

```
In [28]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1460 entries, 0 to 1459
```

```
Data columns (total 63 columns):
```

#	Column	Non-Null Count	Dtype
0	MSSubClass	1460 non-null	int64
1	MSZoning	1460 non-null	object
2	LotArea	1460 non-null	int64
3	LotShape	1460 non-null	object
4	LandContour	1460 non-null	object
5	LotConfig	1460 non-null	object
6	LandSlope	1460 non-null	object
7	Neighborhood	1460 non-null	object
8	Condition1	1460 non-null	object
9	BldgType	1460 non-null	object
10	HouseStyle	1460 non-null	object
11	OverallQual	1460 non-null	int64
12	OverallCond	1460 non-null	int64
13	YearBuilt	1460 non-null	int64
14	YearRemodAdd	1460 non-null	int64
15	RoofStyle	1460 non-null	object
16	Exterior1st	1460 non-null	object
17	Exterior2nd	1460 non-null	object
18	MasVnrArea	1460 non-null	float64
19	ExterQual	1460 non-null	object
20	ExterCond	1460 non-null	object
21	Foundation	1460 non-null	object
22	BsmtQual	1460 non-null	object
23	BsmtCond	1460 non-null	object
24	BsmtExposure	1460 non-null	object
25	BsmtFinType1	1460 non-null	object
26	BsmtFinSF1	1460 non-null	int64
27	BsmtFinType2	1460 non-null	object
28	BsmtFinSF2	1460 non-null	int64
29	BsmtUnfSF	1460 non-null	int64
30	TotalBsmtSF	1460 non-null	int64
31	HeatingQC	1460 non-null	object
32	CentralAir	1460 non-null	object
33	Electrical	1460 non-null	object
34	1stFlrSF	1460 non-null	int64
35	2ndFlrSF	1460 non-null	int64
36	GrLivArea	1460 non-null	int64
37	BsmtFullBath	1460 non-null	int64
38	BsmtHalfBath	1460 non-null	int64
39	FullBath	1460 non-null	int64
40	HalfBath	1460 non-null	int64
41	BedroomAbvGr	1460 non-null	int64
42	KitchenAbvGr	1460 non-null	int64
43	KitchenQual	1460 non-null	object
44	TotRmsAbvGrd	1460 non-null	int64
45	Functional	1460 non-null	object
46	Fireplaces	1460 non-null	int64
47	GarageType	1460 non-null	object
48	GarageYrBlt	1460 non-null	float64
49	GarageFinish	1460 non-null	object
50	GarageCars	1460 non-null	int64
51	GarageArea	1460 non-null	int64
52	GarageQual	1460 non-null	object
53	GarageCond	1460 non-null	object
54	PavedDrive	1460 non-null	object

```

55 WoodDeckSF      1460 non-null int64
56 OpenPorchSF     1460 non-null int64
57 EnclosedPorch   1460 non-null int64
58 MoSold          1460 non-null int64
59 YrSold           1460 non-null int64
60 SaleType        1460 non-null object
61 SaleCondition    1460 non-null object
62 SalePrice       1460 non-null int64

```

dtypes: float64(2), int64(29), object(32)

memory usage: 718.7+ KB

```

In [29]: # Getting categorical variables
cat_var = df.select_dtypes(include='object').columns
cat_var

```

```

Out[29]: Index(['MSZoning', 'LotShape', 'LandContour', 'LotConfig', 'LandSlope',
               'Neighborhood', 'Condition1', 'BldgType', 'HouseStyle', 'RoofStyle',
               'Exterior1st', 'Exterior2nd', 'ExterQual', 'ExterCond', 'Foundation',
               'BsmtQual', 'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinType2',
               'HeatingQC', 'CentralAir', 'Electrical', 'KitchenQual', 'Functional',
               'GarageType', 'GarageFinish', 'GarageQual', 'GarageCond', 'PavedDrive',
               'SaleType', 'SaleCondition'],
              dtype='object')

```

```

In [30]: # Getting numerical variables
num_var = df.select_dtypes(exclude='object').columns
num_var

```

```


Out[30]: Index(['MSSubClass', 'LotArea', 'OverallQual', 'OverallCond', 'YearBuilt',
               'YearRemodAdd', 'MasVnrArea', 'BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF',
               'TotalBsmtSF', '1stFlrSF', '2ndFlrSF', 'GrLivArea', 'BsmtFullBath',
               'BsmtHalfBath', 'FullBath', 'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr',
               'TotRmsAbvGrd', 'Fireplaces', 'GarageYrBlt', 'GarageCars', 'GarageArea',
               'WoodDeckSF', 'OpenPorchSF', 'EnclosedPorch', 'MoSold', 'YrSold',
               'SalePrice'],
              dtype='object')

```

```

In [31]: %matplotlib inline
plt.rcParams['figure.figsize']=6,4
sns.distplot(df['SalePrice'])
plt.show()

```



- The distribution is skewed to the right, as can be seen by looking at it (i.e., outliers on data with high Sales Price). This suggests that there are anomalies present.

```
In [32]: df['SalePrice'].skew()
```

```
Out[32]: 1.8828757597682129
```

The general guideline for determining skewness:

1. In the range of -0.5 to 0.5, the skewness indicates a reasonably symmetric set of data.
2. The data are substantially skewed if the skewness falls between -1 and -0.5 or between 0.5 and 1.
3. The data are significantly skewed if the skewness is greater than 1 or less than -1.

```
In [33]: df['SalePrice'].kurtosis()
```

```
Out[33]: 6.536281860064529
```


- Kurtosis calculates the distribution's tail-heaviness.
- Kurtosis value for a normal distribution is 3.
- The tail is heavier when the kurtosis value rises, and vice versa.

In our case as the kurtosis value is more than ~6.5, distribution tail is heavier

Handling SalePrice high skewness and kurtosis

- To handle this, we will perform Log Transformation on "SalePrice" column. This will transform the variable and make it as normally distributed as possible. Basically it reduces the skewness in the data

```
In [34]: %matplotlib inline
plt.rcParams['figure.figsize']=6,4
sns.displot(np.log(df['SalePrice']), aspect=2)
plt.show()
```



- The data now roughly conforms to a normal distribution.


```
In [35]: df['Transformed_SalePrice'] = np.log(df['SalePrice'])
```

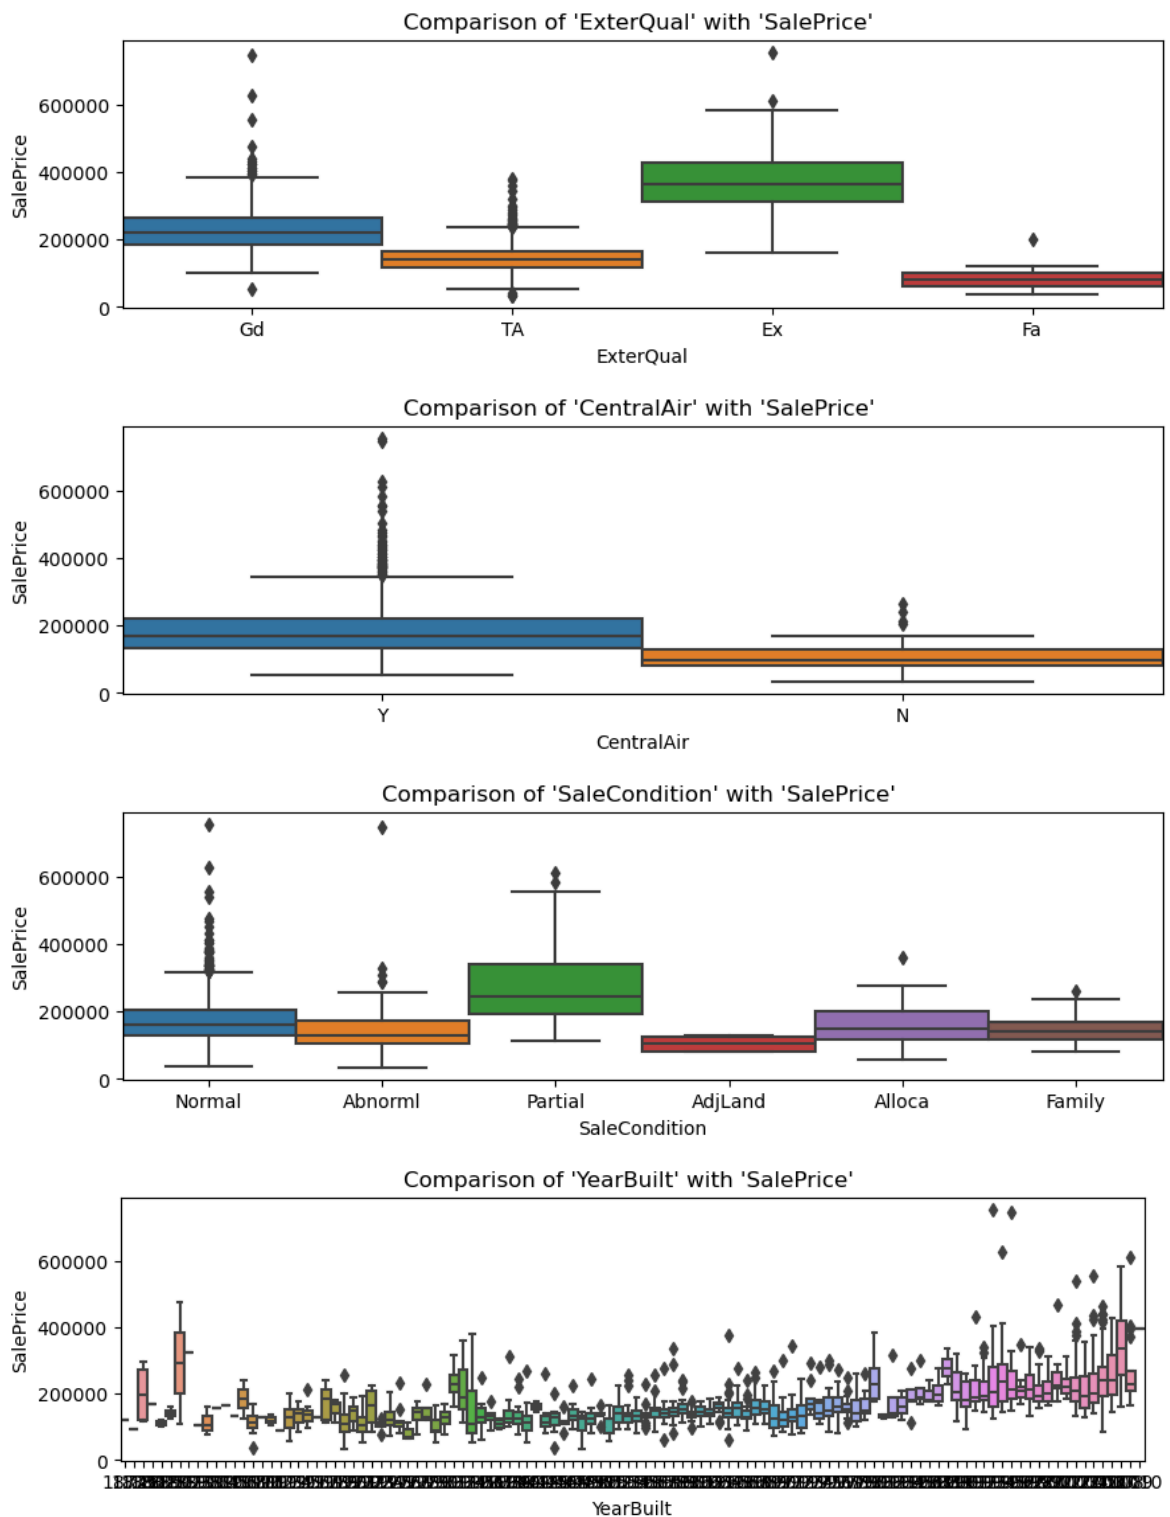
```
In [36]: df.shape
```

```
Out[36]: (1460, 64)
```

EDA- Exploratory Data Analysis

```
In [37]: %matplotlib inline
plt.rcParams['figure.figsize']=10,15
a2=['BsmtQual', 'ExterQual', 'CentralAir', 'SaleCondition', 'YearBuilt']
for i,var in enumerate(a2,start=1):
    plt.subplot(5,1,i)
    sns.boxplot(x=var, y='SalePrice', data=df, width=1)
    plt.title(f"Comparison of '{var}' with 'SalePrice' ")
    plt.show()
```





- As Basement quality increases from Fair to Excellent, we see a corresponding increase in SalePrice
- As Exterior quality increases from Fair to Excellent, we see a corresponding increase in SalePrice
- Houses with Central Air conditioning have a higher median price compared to the houses that don't have Central Air conditioning
- Houses that are partially completed have a higher median Saleprice compared to other categories. This might be because partially completed houses are usually new houses under construction.

- As the house age increases, we can see that the median SalePrice drops but there are few cases where the SalePrice goes up as well

```
In [38]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np

# Assuming df is your DataFrame

# Drop the 'Transformed_SalePrice' column and select only numeric columns
numeric_df = df.drop(columns='Transformed_SalePrice', axis=1).select_dtypes(excl


corr = numeric_df.corr()

# Create a mask for the upper triangle
mask = np.triu(np.ones_like(corr, dtype=bool))

# Set up the matplotlib figure
plt.figure(figsize=[30, 20])

# Draw the heatmap with the mask and correct aspect ratio
sns.heatmap(corr, cmap='RdYlGn', annot=True, square=True, mask=mask)

plt.show()
```




```
In [39]: plt.figure(figsize=[20,10])
```



```

k = 20 # number of variables for a heatmap
cols = corr.nlargest(k, 'SalePrice')['SalePrice'].index
corrmatrix = np.corrcoef(df[cols].values.T)
sns.set(font_scale=1)
hm = sns.heatmap(corrmatrix, cbar=True, annot=True, square=True, fmt='.2f', anno
plt.show()


```



```

In [40]: columns = ['SalePrice', 'OverallQual', 'GrLivArea', 'GarageCars', 'TotalBsmtSF', 'FullBath', 'TotRmsAbvGrd', 'YearBuilt', 'YearRemodAdd', 'MasVnrArea', 'Fireplaces', 'GarageYrBlt', 'BsmtFinSF1', 'WoodDeckSF', '2ndFlrSF', 'OpenPorchSF', 'HalfBath', 'LotArea']
sns.pairplot(df[columns], size=3)
plt.show()

```



```
In [41]: df.shape
```

```
Out[41]: (1460, 64)
```

```
In [42]: from sklearn.preprocessing import LabelEncoder
df_categorical=df.select_dtypes(include='object').columns
lc=LabelEncoder()
for i in df_categorical:
    df[i]=lc.fit_transform(df[i])
```

```
In [43]: df[df_categorical].head()
```

Out[43]:

	MSZoning	LotShape	LandContour	LotConfig	LandSlope	Neighborhood	Condition
0	3	3	3	4	0	5	
1	3	3	3	2	0	24	
2	3	0	3	4	0	5	
3	3	0	3	0	0	6	
4	3	0	3	2	0	15	

In [44]:

df.dtypes

```
Out[44]: MSSubClass      int64
        MSZoning         int32
        LotArea          int64
        LotShape         int32
        LandContour      int32
        LotConfig        int32
        LandSlope        int32
        Neighborhood     int32
        Condition1       int32
        BldgType         int32
        HouseStyle       int32
        OverallQual      int64
        OverallCond      int64
        YearBuilt        int64
        YearRemodAdd     int64
        RoofStyle        int32
        Exterior1st      int32
        Exterior2nd      int32
        MasVnrArea       float64
        ExterQual        int32
        ExterCond        int32
        Foundation       int32
        BsmtQual         int32
        BsmtCond         int32
        BsmtExposure     int32
        BsmtFinType1     int32
        BsmtFinSF1       int64
        BsmtFinType2     int32
        BsmtFinSF2       int64
        BsmtUnfSF        int64
        TotalBsmtSF      int64
        HeatingQC        int32
        CentralAir       int32
        Electrical       int32
        1stFlrSF         int64
        2ndFlrSF         int64
        GrLivArea        int64
        BsmtFullBath     int64
        BsmtHalfBath     int64
        FullBath         int64
        HalfBath         int64
        BedroomAbvGr     int64
        KitchenAbvGr     int64
        KitchenQual      int32
        TotRmsAbvGrd     int64
        Functional       int32
        Fireplaces       int64
        GarageType       int32
        GarageYrBlt      float64
        GarageFinish     int32
        GarageCars       int64
        GarageArea       int64
        GarageQual       int32
        GarageCond       int32
        PavedDrive       int32
        WoodDeckSF       int64
        OpenPorchSF      int64
        EnclosedPorch    int64
        MoSold           int64
        YrSold           int64
```

```

SaleType          int32
SaleCondition      int32
SalePrice          int64
Transformed_SalePrice float64
dtype: object

```

In [45]: `df.shape`

Out[45]: (1460, 64)

```

In [46]: #Divide data into X and y for building the model
x=df.drop(['SalePrice','Transformed_SalePrice'], axis=1)
y=df['Transformed_SalePrice']

```

In [47]: `x.head()`

Out[47]:

	MSSubClass	MSZoning	LotArea	LotShape	LandContour	LotConfig	LandSlope	Ne
0	60	3	8450	3	3	4	0	
1	20	3	9600	3	3	2	0	
2	60	3	11250	0	3	4	0	
3	70	3	9550	0	3	0	0	
4	60	3	14260	0	3	2	0	

In [48]: `y.head()`

Out[48]:

```

0    12.247694
1    12.109011
2    12.317167
3    11.849398
4    12.429216
Name: Transformed_SalePrice, dtype: float64

```

```

In [49]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=0)

```

```

In [50]: from sklearn.preprocessing import StandardScaler
sc=StandardScaler()
x_train=sc.fit_transform(x_train)
x_test=sc.fit_transform(x_test)

```

Building base Model

```

In [51]: from sklearn.linear_model import (
LinearRegression,Ridge,Lasso,ElasticNet,SGDRegressor,HuberRegressor
)
from sklearn.neighbors import KNeighborsRegressor
from sklearn.svm import SVR
from sklearn.ensemble import RandomForestRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.preprocessing import PolynomialFeatures

```

```

from sklearn.pipeline import Pipeline
from sklearn.neural_network import MLPRegressor
import lightgbm as lgb
import xgboost as xgb
from sklearn.metrics import r2_score, mean_squared_error, mean_absolute_error

```

```

In [52]: #Define the models
models={
    'Linear Regression':LinearRegression(),
    'Robust Regression':HuberRegressor(),
    'Ridge Regression':Ridge(),
    'ElasticNet Regressor': ElasticNet(),
    'Lasso Regression':Lasso(),
    'Polynomial Regression':Pipeline([
        ('poly',PolynomialFeatures(degree=2)),
        ('linear',LinearRegression())
    ]),
    'SGD Regressor':SGDRegressor(),
    'ANN':MLPRegressor(hidden_layer_sizes=(100,),max_iter=1000),
    'Random Forest Regressor':RandomForestRegressor(),
    'Support vector Regressor':SVR(),
    'LGBM':lgb.LGBMRegressor(),
    'XGBoost':xgb.XGBRFRegressor(),
    'KNN Regressor':KNeighborsRegressor()
}
results=[]

```

```

In [ ]: for name,model in models.items():
    model.fit(x_train,y_train)
    y_pred=model.predict(x_test)
    y_pred_train=model.predict(x_train)
    r2_train = r2_score(y_train, y_pred_train)
    r2_test=r2_score(y_test, y_pred)
    mse_train = mean_squared_error(y_train, y_pred_train)
    mse_test=mean_squared_error(y_test, y_pred)
    mae_train=mean_absolute_error(y_train, y_pred_train)
    mae_test=mean_absolute_error(y_test, y_pred)
    results.append({
        'Name of the model':name,
        'r2_train':r2_train,
        'r2_test':r2_test,
        'mse_train':mse_train,
        'mse_test':mse_test,
        'mae_train':mae_train,
        'mae_test':mae_test
    })
df1=pd.DataFrame(results)

```

```

In [ ]: df1

```

To determine which model is the best, we should consider the balance between training and testing performance, the complexity of the model, and the overall error metrics. Here are the considerations for each model:

Key Metrics for Evaluation:

- **R² Score:** Indicates the proportion of variance explained by the model. Higher is generally better, but it should be similar for both training and test sets to avoid overfitting.
- **MSE (Mean Squared Error):** Measures the average squared difference between predicted and actual values. Lower is better.
- **MAE (Mean Absolute Error):** Measures the average absolute difference between predicted and actual values. Lower is better.

Model Performance Summary:

1. Linear Regression:

- **R² Train:** 0.898227
- **R² Test:** -1.153681e+22 (significantly negative)
- **MSE Train:** 0.016442
- **MSE Test:** 1.782565e+21
- **MAE Train:** 0.086249
- **MAE Test:** 3.276174e+10
- **Conclusion:** Severe overfitting, poor generalization. Not a good model.

2. KNeighbor Regressor:

- **R² Train:** 0.869135
- **R² Test:** 0.793835
- **MSE Train:** 0.021143
- **MSE Test:** 0.031849
- **MAE Train:** 0.101664
- **MAE Test:** 0.128370
- **Conclusion:** Balanced performance, good generalization. A good model.

3. Support Vector Regressor:

- **R² Train:** 0.965355
- **R² Test:** 0.864931
- **MSE Train:** 0.005597
- **MSE Test:** 0.020870
- **MAE Train:** 0.063692
- **MAE Test:** 0.098854
- **Conclusion:** Slight overfitting, but overall good performance. A good model.

4. Random Forest Regressor:

- **R² Train:** 0.980918
- **R² Test:** 0.889454
- **MSE Train:** 0.003083
- **MSE Test:** 0.017081
- **MAE Train:** 0.036678
- **MAE Test:** 0.089254
- **Conclusion:** Slight overfitting, but excellent performance. A good model.

5. Decision Tree Regressor:

- **R² Train:** 1.000000
- **R² Test:** 0.750913
- **MSE Train:** 0.000000
- **MSE Test:** 0.038487
- **MAE Train:** 0.000000
- **MAE Test:** 0.141652
- **Conclusion:** Severe overfitting, poor generalization. Not a good model.

Best Model Selection:

- **Random Forest Regressor:** It shows the highest R² on the test set, indicating the best generalization performance. Although there is slight overfitting, it is within acceptable limits, and the overall error metrics are low.
- **Support Vector Regressor:** Also a good option with balanced performance, though slightly lower than the Random Forest.

Conclusion:

The Random Forest Regressor is the best model based on the balance between training and test performance, as well as low error metrics. It shows excellent generalization and maintains low errors, making it a robust choice for your problem.

```
In [ ]: x.columns
```

```
In [ ]: # to build the front end i need to identify the key attributes.
model = RandomForestRegressor(n_estimators=100, random_state=0)
model.fit(x_train, y_train)
importances = model.feature_importances_
features = x.columns
feature_importance_df = pd.DataFrame({'Feature': features, 'Importance': importances})
feature_importance_df = feature_importance_df.sort_values(by='Importance', ascending=False)
```

```
In [ ]: feature_importance_df
```

```
In [ ]: feature_importance_df.to_csv('feature_importance.csv', index=False)
```